

Day 3: Programming with Python

Pandas Basics

Pandas is one of the most used open-source Python libraries to work with structured tabular data for analysis. Pandas is widely used for data science, machine learning, and many more. It offers data structures and operations for manipulating numerical tables and time series.

In this unit we will see some of the basic functionalities of Pandas like how we can create Pandas data structures, view information on the data and columns, working with index, and how to combine different pandas data structures.

0. Import Pandas

The Pandas website where you can find everything about Pandas is here:

<https://pandas.pydata.org/>.

We use `import pandas as pd` to import Pandas. The alias `pd` is commonly used for Pandas.

```
In [1]: import pandas as pd
```

1. Pandas datatypes

Key points:

- **`pd.Series()`** - Creates a Series
- **`pd.DataFrame()`** - Creates a DataFrame

There are two main data structures in Pandas, `Series` and `DataFrames`.

`Series` is a one-dimensional data structure. It holds data of many types including objects, floats, strings and integers. `DataFrame` is a more complex data structure and is designed to contain data with several dimensions.

1.1 Series

To create a Series we call the `Series()` function and pass as an argument an array containing the values to be included in it.

Let's create our first Series from a list (`listItems`) that contains the numbers `[1,3,6,9]`.

```
In [2]: listItems = [1,3,6,9]
s = pd.Series(listItems)
print (s)
```

```
1    3
2    6
3    9
dtype: int64
```

The Series has a `dtype`, that refers to the type of the data.

In the above example, we see that the `dtype` is integer. This is because our list contained only integers.

Let's now try to create a Series from a list that contains the numbers [1.1,3.1,6.1,9.1] and let's check if the type of the data has changed. What do you notice?

```
In [63]: listItems = [1.1,3.1,6.1,9.1]
s = pd.Series(listItems)
print(s)
```

```
float64
```

Let's create a Series that contains ["Anna", "John", "Mark"]. What is going to be the type of the Series?

```
In [4]: listItems = ["Anna", "John", "Mark"]
s = pd.Series(listItems)
print(s)
```

```
0    Anna
1    John
2    Mark
dtype: object
```

1.2 Pandas DataFrames

The `DataFrame` is a tabular data structure and is similar to a spreadsheet. It is a 2-dimensional data structure where each column contains data of the same type. DataFrames are great for representing real data: rows correspond to instances, and columns correspond to features of these instances.

For example, in case of movies, we have movies as rows and different metadata (title, release year, duration etc.) for each movie in the columns.

MovieID	title	year	duration
0	The Pianist	2002	150
1	Gladiator	2000	155
2	The Godfather	1972	177
3	Inception	2010	148
4	Titanic	1997	195

Create DataFrame from existing data structure

There are different ways to create a DataFrame. It can be created by typing the values, or from an existing data structure (list, dictionary) or imported from a file.

To create a DataFrame we use the `pd.DataFrame()` function.

Let's create a simple DataFrame from a single list, [1,3,6,9]

```
In [5]: listItems = [1,3,6,9]
        pd.DataFrame(listItems)
```

```
Out[5]: 0
        0  1
        1  3
        2  6
        3  9
```

Let's create a second DataFrame from a dictionary. The dictionary contains the title, release year and duration in minutes of some movies.

```
In [6]: data = {'title': ['The Pianist', 'Gladiator', 'The Godfather', 'Inception',
                          'Titanic'],
                'year': [2002, 2000, 1972, 2010, 1997],
                'duration': [150, 155, 177, 148, 195]}

data
```

```
Out[6]: {'title': ['The Pianist',
                  'Gladiator',
                  'The Godfather',
                  'Inception',
                  'Titanic'],
         'year': [2002, 2000, 1972, 2010, 1997],
         'duration': [150, 155, 177, 148, 195]}
```

Now we create the DataFrame from the dictionary

```
In [7]: df = pd.DataFrame(data)
        print(df)
```

	title	year	duration
0	The Pianist	2002	150
1	Gladiator	2000	155
2	The Godfather	1972	177
3	Inception	2010	148
4	Titanic	1997	195

2. Index

Key points:

- **pd.Series(list, index)** - Creates a Series and assigns index
- **series.index = list** - Assigns index to Series
- **pd.DataFrame(dict, index)** - Creates a DataFrame and assigns index
- **df.index** - Returns the index of the DataFrame
- **df.values** - Returns the values of the DataFrame

Until now, every time we created a Series from a list there was an additional column on the left. This axis is called `index` and is added into the Series and DataFrames. Index is like an address for the stored data, and can be used to access any data point across the Series or DataFrame.

The index can be specified at the creation of the Series or the DataFrame. If we do not specify any index during the construction of the Pandas object then, by default, Pandas will assign numerical values increasing from 0 as labels.

```
In [8]: s = pd.Series(listItems)
s
```

```
Out[8]: 0    1
        1    3
        2    6
        3    9
        dtype: int64
```

Let's create a Series that contains grades of a student (listItems) regarding the following 4 courses ['Math', 'Physics', 'Chemistry', 'History']. We can set the courses as index with the parameter `index`

```
In [9]: s = pd.Series(listItems, index = ['Math', 'Physics', 'Chemistry', 'History'])
print(s)
```

```
Math          1
Physics        3
Chemistry      6
History        9
dtype: int64
```

The index can also be defined later with the `s.index` attribute.

```
In [10]: s = pd.Series(listItems)
s.index = ['Math', 'Physics', 'Chemistry', 'History']
print(s)
```

```
Math          1
Physics        3
Chemistry      6
History        9
dtype: int64
```

In a similar way we define the index in DataFrames.

```
In [11]: movies = ['The Pianist', 'Gladiator', 'The Godfather', 'Inception',
                  'Titanic', 'Moulin Rouge', 'La La Land', 'The Notebook']
df = pd.DataFrame({'year': [2002, 2000, 1972, 2010, 1997, 2001, 2016, 2004],
                  'duration': [150, 155, 177, 148, 195, 130, 128, 124]},
                  index=movies)

print(df)
```

	year	duration
The Pianist	2002	150
Gladiator	2000	155
The Godfather	1972	177
Inception	2010	148
Titanic	1997	195
Moulin Rouge	2001	130
La La Land	2016	128
The Notebook	2004	124

We can use the `df.values` and the `df.index` attributes to view the value and index arrays

```
In [12]:
```

```
df.values
```

```
Out[12]: array([[2002, 150],
               [2000, 155],
               [1972, 177],
               [2010, 148],
               [1997, 195],
               [2001, 130],
               [2016, 128],
               [2004, 124]])
```

```
In [13]: df.index
```

```
Out[13]: Index(['The Pianist', 'Gladiator', 'The Godfather', 'Inception', 'Titanic',
               'Moulin Rouge', 'La La Land', 'The Notebook'],
               dtype='object')
```

3. Knowing my data

Key points:

- **df.info()** - Prints a concise summary of a DataFrame
- **df.shape** - Returns a tuple representing the dimensionality of the DataFrame
- **df.columns** - Returns the column labels of the DataFrame
- **df.head(n)** - Returns the first n rows (default is 5)
- **df.tail(n)** - Returns the last n rows (default is 5)

Let's see how we can know more about the data

```
In [14]: df = pd.DataFrame({'year':[2002, 2000, 1972, 2010, 1997, 2001, 2016, 2004],
                             'duration':[150, 155, 177, 148, 195, 130, 128, 124]},
                             index=movies)

print(df)
```

	year	duration
The Pianist	2002	150
Gladiator	2000	155
The Godfather	1972	177
Inception	2010	148
Titanic	1997	195
Moulin Rouge	2001	130
La La Land	2016	128
The Notebook	2004	124

There are some functions to get some basic information regarding our DataFrame.

We use the `info()` to output some general information about the DataFrame, the `shape` to get its dimensions, and the `columns` attribute to get the column labels of the DataFrame. Note that `shape` and `columns` are attributes and not functions and that is why they do not have parentheses.

With the `info()` we can see the columns of the DataFrame, the number of non-null values and the type of the data per column.

```
In [15]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 8 entries, The Pianist to The Notebook
```

```
Data columns (total 2 columns):
#   Column   Non-Null Count  Dtype
---  -
0   year     8 non-null        int64
1   duration  8 non-null        int64
dtypes: int64(2)
memory usage: 192.0+ bytes
```

With the `shape` we get the dimensions of the DataFrame

```
In [16]: df.shape
```

```
Out[16]: (8, 2)
```

With the `columns` we get the names of the columns

```
In [17]: df.columns
```

```
Out[17]: Index(['year', 'duration'], dtype='object')
```

We can get the first rows of the DataFrame with the `head()` function. The `head()` function shows the first five rows if there is no argument or will show the number of rows specified in the argument. The function `tail()` shows the last entries of the DataFrame.

```
In [18]: df.head()
```

```
Out[18]:
```

	year	duration
The Pianist	2002	150
Gladiator	2000	155
The Godfather	1972	177
Inception	2010	148
Titanic	1997	195

Viewing the first 3 rows

```
In [19]: df.head(3)
```

```
Out[19]:
```

	year	duration
The Pianist	2002	150
Gladiator	2000	155
The Godfather	1972	177

To view the last rows we use the `tail()` function

```
In [20]: df.tail()
```

```
Out[20]:
```

	year	duration
Inception	2010	148
Titanic	1997	195

	year	duration
Moulin Rouge	2001	130
La La Land	2016	128
The Notebook	2004	124

4. Working with columns

Key points:

- **df.column** - Returns the data of a column
- **df['column']** - Returns the data of a column
- **df[['column1', 'column2']]** - Returns the data of multiple columns
- **df['new_column'] = list** - Creates a new column and assigns data to it
- **del df['column']** - Deletes a column

For now we will work with only the DataFrame we created before.

```
In [21]: df = pd.DataFrame({'year':[2002, 2000, 1972, 2010, 1997, 2001, 2016, 2004],
                           'duration':[150, 155, 177, 148, 195, 130, 128, 124]},
                           index=movies)

print(df)
```

	year	duration
The Pianist	2002	150
Gladiator	2000	155
The Godfather	1972	177
Inception	2010	148
Titanic	1997	195
Moulin Rouge	2001	130
La La Land	2016	128
The Notebook	2004	124

To select a column, we use the column name.

```
In [22]: df.year
```

```
Out[22]: The Pianist      2002
Gladiator      2000
The Godfather   1972
Inception      2010
Titanic        1997
Moulin Rouge   2001
La La Land     2016
The Notebook    2004
Name: year, dtype: int64
```

```
In [23]: df['year']
```

```
Out[23]: The Pianist      2002
Gladiator      2000
The Godfather   1972
Inception      2010
Titanic        1997
Moulin Rouge   2001
La La Land     2016
```

```
The Notebook      2004
Name: year, dtype: int64
```

If we want data of more columns we provide a list of the column names

```
In [24]: df[['year', 'duration']]
```

```
Out[24]:
```

	year	duration
The Pianist	2002	150
Gladiator	2000	155
The Godfather	1972	177
Inception	2010	148
Titanic	1997	195
Moulin Rouge	2001	130
La La Land	2016	128
The Notebook	2004	124

Add a column

Let's say we want to add the column of `genre` in our DataFrame and we know that for the movies the genres are ['drama','action','crime','action','romance']

```
In [25]: genre = ['drama', 'action', 'crime', 'action', 'romance',
                  'drama', 'romance', 'romance']
df['genre'] = genre
df
```

```
Out[25]:
```

	year	duration	genre
The Pianist	2002	150	drama
Gladiator	2000	155	action
The Godfather	1972	177	crime
Inception	2010	148	action
Titanic	1997	195	romance
Moulin Rouge	2001	130	drama
La La Land	2016	128	romance
The Notebook	2004	124	romance

Delete a column

To delete a column, we use the keyword `del` followed by the column to be deleted (e.g., `del df['genre']`)

```
In [26]: del df['genre']
df
```

```
Out[26]:
```

	year	duration
--	------	----------

	year	duration
The Pianist	2002	150
Gladiator	2000	155
The Godfather	1972	177
Inception	2010	148
Titanic	1997	195
Moulin Rouge	2001	130
La La Land	2016	128
The Notebook	2004	124

5. Working with rows

Key points:

- **df.loc['index']** - Returns data of the row with index label index
- **df.loc[['index1', 'index2']]** - Returns data of multiple rows
- **df.loc['index1': 'index2']** - Returns data from row index1 until index2
- **df.loc['index1': 'index2', 'column']** - Returns data of a column and from row index1 until index2
- **df.loc[df['column'] > /< />= /<= /==]** - Returns data of a column based on a condition
- **df.iloc[indexID]** - Returns data of a row with specific index id

5.1. Select a row with loc

We can select a row using the `.loc` that can access group of values using labels. For example the code `df.loc['Gladiator']` will return the data of the row of `Gladiator`

```
In [27]: df.loc['Gladiator']
```

```
Out[27]: year      2000
duration    155
Name: Gladiator, dtype: int64
```

Let's say that now we want to return the data for `Gladiator` and `Titanic`. The command `df.loc['Gladiator', 'Titanic']` will return an error.

To return more rows we use a list of labels. Note that using `[[]]` returns a DataFrame.

```
In [28]: df.loc[['Gladiator', 'Titanic']]
```

```
Out[28]:
```

	year	duration
Gladiator	2000	155
Titanic	1997	195

We can also access data with the slice. Let's return the data from `Gladiator` until `Inception`

```
In [29]:
```

```
df.loc['Gladiator':'Inception']
```

```
Out[29]:
```

	year	duration
Gladiator	2000	155
The Godfather	1972	177
Inception	2010	148

Also, we can select to get only the values of a specific column. The first argument is the rows and the second the columns that we want to access. So we want only the `duration` of the movies from `Gladiator` until `Inception`

```
In [30]: df.loc['Gladiator':'Inception', 'duration']
```

```
Out[30]:
```

Gladiator	155
The Godfather	177
Inception	148

Name: duration, dtype: int64

5.2 Conditional selection

In different cases we want to select data that meet a specific condition. For example, we want to get the data of movies that last more that 2.5 hours. We can first set the expression `df['duration'] > 150`. This will return a boolean value for every row

```
In [31]: df['duration'] > 150
```

```
Out[31]:
```

The Pianist	False
Gladiator	True
The Godfather	True
Inception	False
Titanic	True
Moulin Rouge	False
La La Land	False
The Notebook	False

Name: duration, dtype: bool

To get the rows we can pass the expression in the `loc` attribute. Only the rows with that were assigned as `True` will be printed

```
In [32]: df.loc[df['duration'] > 150]
```

```
Out[32]:
```

	year	duration
Gladiator	2000	155
The Godfather	1972	177
Titanic	1997	195

5.3 Select a row with iloc

The `iloc` selects data only by integer index.

As we mentioned before every row and column in a DataFrame has an integer location that is assigned to it. This is even if we have assigned our own index. The integer location is simply

the number of rows/columns from the top/left beginning at 0.

To return the third row we use `df.iloc[2]`

```
In [33]: df.iloc[2]
```

```
Out[33]: year      1972
duration    177
Name: The Godfather, dtype: int64
```

6. Combining DataFrames

Key points:

- **pd.concat([df1, df2])** - Concatenates two DataFrames
- **pd.merge(df1, df2)** - Merges two DataFrames similar to relational algebra

Pandas provides various facilities for combining together Series or DataFrames with various kinds of set logic for the indexes and relational algebra functionality in the case of merge-type operations.

6.1 Concatenate two DataFrames

The `concat()` function does all of the heavy lifting of performing concatenation operations along an axis while performing optional set logic (union or intersection) of the indexes (if any) on the other axes.

Here is a simple example. Let's say we have our DataFrame from before with movies.

```
In [34]: df = pd.DataFrame({'year':[2002, 2000, 1972, 2010, 1997],
                           'duration':[150, 155, 177, 148, 195]},
                           index=['The Pianist', 'Gladiator', 'The Godfather',
                                  'Inception', 'Titanic'])

df
```

```
Out[34]:
```

	year	duration
The Pianist	2002	150
Gladiator	2000	155
The Godfather	1972	177
Inception	2010	148
Titanic	1997	195

We also have data on more movies which maybe they came from a different source and we want to combine these two DataFrames. We have Nomadland (2021) that lasts 108 minutes and Parasite (2019) that lasts 132 minutes

```
In [35]: df_new = pd.DataFrame({'year': [2021, 2019],
                                'duration': [108, 132]},
                                index=['Nomadland', 'Parasite'])

df_new
```

```
Out[35]:
```

	year	duration
Nomadland	2021	108
Parasite	2019	132

In order to concatenate the two DataFrames, we call the `concat()` function. The `concat()` function takes as argument the list of the DataFrames.

```
In [36]: df_conc = pd.concat([df, df_new])
df_conc
```

```
Out[36]:
```

	year	duration
The Pianist	2002	150
Gladiator	2000	155
The Godfather	1972	177
Inception	2010	148
Titanic	1997	195
Nomadland	2021	108
Parasite	2019	132

Let's say we get data on the genres of the movies. We can concatenate the DataFrames again but not we have to set the axis to 1

```
In [37]: df_new = pd.DataFrame({'genre': ['drama', 'action', 'crime', 'action', 'romance'],
                                index=['The Pianist', 'Gladiator', 'The Godfather',
                                        'Inception', 'Titanic']})

df_new
```

```
Out[37]:
```

	genre
The Pianist	drama
Gladiator	action
The Godfather	crime
Inception	action
Titanic	romance

```
In [38]: df_conc = pd.concat([df, df_new])
df_conc
```

```
Out[38]:
```

	year	duration	genre
The Pianist	2002.0	150.0	NaN
Gladiator	2000.0	155.0	NaN
The Godfather	1972.0	177.0	NaN
Inception	2010.0	148.0	NaN
Titanic	1997.0	195.0	NaN

	year	duration	genre
The Pianist	NaN	NaN	drama
Gladiator	NaN	NaN	action
The Godfather	NaN	NaN	crime
Inception	NaN	NaN	action
Titanic	NaN	NaN	romance

```
In [39]: df_conc = pd.concat([df, df_new], axis = 1)
df_conc
```

```
Out[39]:
```

	year	duration	genre
The Pianist	2002	150	drama
Gladiator	2000	155	action
The Godfather	1972	177	crime
Inception	2010	148	action
Titanic	1997	195	romance

6.2 Merge two DataFrames

We use the `merge()` function to combine data objects based on one or more keys in a similar way to a relational database. More specifically, `merge()` is most useful when we want to combine rows that share data.

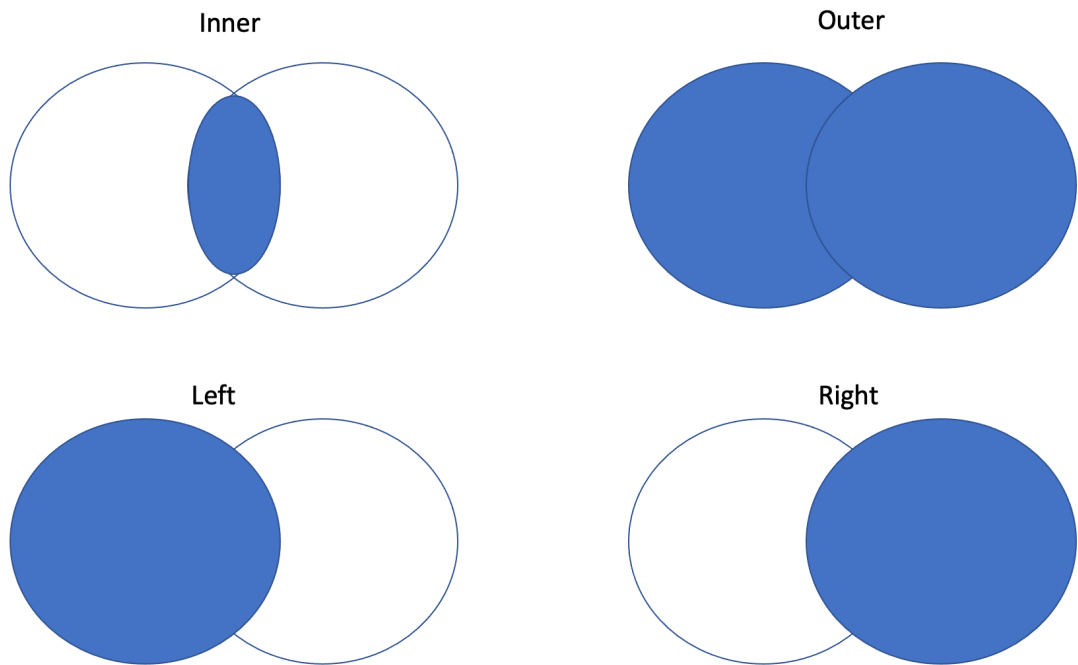
We can achieve both many-to-one and many-to-many joins with `merge()`. In a many-to-one join, one of your DataFrames will have many rows in the merge column that repeat the same values, while the merge column in the other DataFrame will not have repeat values.

In a many-to-many join, both of the merge columns have repeat values. These merges are more complex. This means that, after the merge, we have every combination of rows that share the same value in the key column.

When we use `merge()`, we provide two required arguments: the left DataFrame and the right DataFrame.

After that, we can provide a number of optional arguments to define how the datasets are merged, such as `how` that defines the kind of merge.

![howMerge.png]



Let's see some examples:

The simplest type of merge is the one-to-one join. Let's continue with our examples from above with some modifications. Let's create the `df` as:

```
In [40]: df = pd.DataFrame({'title': ['The Pianist', 'Gladiator', 'The Godfather',  
                                     'Inception', 'Titanic'],  
                           'year': [2002, 2000, 1972, 2010, 1997],  
                           'duration': [150, 155, 177, 148, 195]})  
  
print(df)
```

	title	year	duration
0	The Pianist	2002	150
1	Gladiator	2000	155
2	The Godfather	1972	177
3	Inception	2010	148
4	Titanic	1997	195

And now we know that *Gladiator* has a score of 8.9 and *Inception* of 8.5. We first create the `df_new` DataFrame

```
In [41]: df_new = pd.DataFrame({'title': ['Gladiator', 'Inception'],  
                               'score': [8.9, 8.5]})  
  
print(df_new)
```

	title	score
0	Gladiator	8.9
1	Inception	8.5

If we try to concatenate the two DataFrames, we will have the following result

```
In [42]: df_concat = pd.concat([df, df_new])  
df_concat
```

```
Out[42]:
```

	title	year	duration	score
--	-------	------	----------	-------

	title	year	duration	score
0	The Pianist	2002.0	150.0	NaN
1	Gladiator	2000.0	155.0	NaN
2	The Godfather	1972.0	177.0	NaN
3	Inception	2010.0	148.0	NaN
4	Titanic	1997.0	195.0	NaN
0	Gladiator	NaN	NaN	8.9
1	Inception	NaN	NaN	8.5

So the `concat()` method ignores the title and the fact that there are already data for Gladiator and Inception

We have to use `merge()` in this case:

In [43]:

```
df_merge = pd.merge(df, df_new)
print(df_merge)
print()
```

	title	year	duration	score
0	Gladiator	2000	155	8.9
1	Inception	2010	148	8.5

The `pd.merge()` function recognizes that each DataFrame has a `title` column, and automatically joins using this column as a key. The result of the merge is a new DataFrame that combines the information from the two input DataFrames.

The order of entries in each column is not necessarily maintained: in this case, the order of the title column is different between `df` and `df_new`.

Many-to-one joins

Many-to-one joins are joins in which one of the two key columns contains duplicate entries. For the many-to-one case, the resulting DataFrame preserves those duplicate entries. Consider the following example of a many-to-one join:

In [44]:

```
df_movies = pd.DataFrame({'title': ['Gladiator', 'Amelie', 'Parasite'],
                           'duration': [155, 123, 132]})
df_genre = pd.DataFrame({'title': ['Gladiator', 'Gladiator', 'Amelie',
                                    'Titanic'],
                           'genre': ['action', 'drama', 'romance', 'romance']})

print(df_movies)
print()

print(df_genre)
print()

print(pd.merge(df_movies, df_genre))
```

	title	duration
0	Gladiator	155
1	Amelie	123
2	Parasite	132

	title	genre
0	Gladiator	action

1	Gladiator	drama
2	Amelie	romance
3	Titanic	romance

	title	duration	genre
0	Gladiator	155	action
1	Gladiator	155	drama
2	Amelie	123	romance

We have many-to-many merge if the key column in both the left and right array contains duplicates.

In the following example, we have a DataFrame showing one or more countries related to the movie and a DataFrame showing the genres of the movies. By performing a many-to-many join, we associate movies with the genres and countries:

```
In [45]: df_country = pd.DataFrame({'title': ['Gladiator', 'Gladiator', 'Amelie',
                                             'Amelie', 'Parasite'],
                                   'country': ['UK', 'USA', 'France', 'Germany',
                                              'South Korea']})

df_genre = pd.DataFrame({'title': ['Gladiator', 'Gladiator', 'Amelie',
                                    'Titanic'],
                         'genre': ['action', 'drama', 'romance',
                                   'romance']})

print(df_country)
print()

print(df_genre)
print()

print(pd.merge(df_genre, df_country))
```

	title	country
0	Gladiator	UK
1	Gladiator	USA
2	Amelie	France
3	Amelie	Germany
4	Parasite	South Korea

	title	genre
0	Gladiator	action
1	Gladiator	drama
2	Amelie	romance
3	Titanic	romance

	title	genre	country
0	Gladiator	action	UK
1	Gladiator	action	USA
2	Gladiator	drama	UK
3	Gladiator	drama	USA
4	Amelie	romance	France
5	Amelie	romance	Germany

The how argument

The `how` argument defines the type of merge to be performed. This argument is by default to `inner`. Let's say we have the `df_country` and `df_genre` from before

Let's see the following example:

```
In [46]: print(df_country)
print()
```



```
print(df_genre)
print()
```

	title	country
0	Gladiator	UK
1	Gladiator	USA
2	Amelie	France
3	Amelie	Germany
4	Parasite	South Korea

	title	genre
0	Gladiator	action
1	Gladiator	drama
2	Amelie	romance
3	Titanic	romance

Let's compare the inner and outer intersections.

In [47]:

```
print(pd.merge(df_country, df_genre, how='inner'))
print()

print(pd.merge(df_country, df_genre, how='outer'))
print()
```

	title	country	genre
0	Gladiator	UK	action
1	Gladiator	UK	drama
2	Gladiator	USA	action
3	Gladiator	USA	drama
4	Amelie	France	romance
5	Amelie	Germany	romance

	title	country	genre
0	Gladiator	UK	action
1	Gladiator	UK	drama
2	Gladiator	USA	action
3	Gladiator	USA	drama
4	Amelie	France	romance
5	Amelie	Germany	romance
6	Parasite	South Korea	NaN
7	Titanic	NaN	romance

Summary

In this unit, we covered basic functions of Pandas including how to create Pandas Series and DataFrames, and how to access the data.

In addition, we covered the main joining functions of Pandas, namely `concat()` and `merge()`.

Although these functions operate quite similar to each other, there are some fundamental differences among them.

Pandas `concat()` can be used for joining multiple DataFrames through both columns or rows. It is considered to be the most efficient method of joining DataFrames.

`Merge()` function performs joins similar to SQL. With the help of `merge()` we can merge values using a common column found in two DataFrames.

Exercises

1. Read the IMDB_movies.csv file and print all the basic information for the data, names of columns, shape and the first 5 rows

In [48]:

```
movies = pd.read_csv("../data/IMDB_movies.csv")

print(movies.columns)
print(movies.info())
print(movies.index)

print(movies.head())
```

```
Index(['imdb_title_id', 'title', 'original_title', 'year', 'date_published',
      'genre', 'duration', 'country', 'language', 'director', 'writer',
      'production_company', 'actors', 'description', 'avg_vote', 'votes',
      'budget', 'usa_gross_income', 'worldwide_gross_income', 'metascore',
      'reviews_from_users', 'reviews_from_critics'],
      dtype='object')
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999
```

```
Data columns (total 22 columns):
```

#	Column	Non-Null Count	Dtype
0	imdb_title_id	10000 non-null	object
1	title	10000 non-null	object
2	original_title	10000 non-null	object
3	year	10000 non-null	int64
4	date_published	10000 non-null	object
5	genre	10000 non-null	object
6	duration	10000 non-null	int64
7	country	10000 non-null	object
8	language	9803 non-null	object
9	director	9999 non-null	object
10	writer	9966 non-null	object
11	production_company	9971 non-null	object
12	actors	9999 non-null	object
13	description	9921 non-null	object
14	avg_vote	10000 non-null	float64
15	votes	10000 non-null	int64
16	budget	1880 non-null	object
17	usa_gross_income	451 non-null	object
18	worldwide_gross_income	392 non-null	object
19	metascore	227 non-null	float64
20	reviews_from_users	9821 non-null	float64
21	reviews_from_critics	9236 non-null	float64

```
dtypes: float64(4), int64(3), object(15)
```

```
memory usage: 1.7+ MB
```

```
None
```

```
RangeIndex(start=0, stop=10000, step=1)
```

	imdb_title_id	title	original_title	\
0	tt0000009	Miss Jerry	Miss Jerry	
1	tt0000574	The Story of the Kelly Gang	The Story of the Kelly Gang	
2	tt0001892	Den sorte drøm	Den sorte drøm	
3	tt0002101	Cleopatra	Cleopatra	
4	tt0002130	L'Inferno	L'Inferno	

	year	date_published	genre	duration	country
0	1894	1894-10-09	Romance	45	USA
1	1906	1906-12-26	Biography, Crime, Drama	70	Australia
2	1911	1911-08-19	Drama	53	Germany, Denmark
3	1912	1912-11-13	Drama, History	100	USA
4	1911	1911-03-06	Adventure, Drama, Fantasy	68	Italy

```

    language                                director ... \
0      None                                Alexander Black ...
1      None                                Charles Tait ...
2      NaN                                Urban Gad ...
3  English                                Charles L. Gaskill ...
4  Italian  Francesco Bertolini, Adolfo Padovan ...

                                actors \
0  Blanche Bayliss, William Courtenay, Chauncey D...
1  Elizabeth Tait, John Tait, Norman Campbell, Be...
2  Asta Nielsen, Valdemar Psilander, Gunnar Helse...
3  Helen Gardner, Pearl Sindelar, Miss Fielding, ...
4  Salvatore Papa, Arturo Pirovano, Giuseppe de L...

                                description avg_vote votes    budget
\
0  The adventures of a female reporter in the 1890s.      5.9   154      NaN
1  True story of notorious Australian outlaw Ned ...     6.1   589    $ 2250
2  Two men of high rank are both wooing the beaut...     5.8   188      NaN
3  The fabled queen of Egypt's affair with Roman ...     5.2   446    $ 45000
4  Loosely adapted from Dante's Divine Comedy and...     7.0  2237      NaN

    usa_gross_income worldwide_gross_income metascore reviews_from_users \
0              NaN              NaN              NaN              1.0
1              NaN              NaN              NaN              7.0
2              NaN              NaN              NaN              5.0
3              NaN              NaN              NaN             25.0
4              NaN              NaN              NaN             31.0

    reviews_from_critics
0              2.0
1              7.0
2              2.0
3              3.0
4             14.0

[5 rows x 22 columns]

```

2. Assume you have the following dictionary with some passengers of a flight. Create a DataFrame from the dictionary and assign pass_id as index. Print the basic information on the DataFrame and the index and the first 7 rows

```
In [49]: passengers = {"age":[23, 25, 78, 12, 56, 33, 67, 78, 34, 64], "priority":["y'
pass_id=[101,102,103,104,105,106,107,108,109,110]
```

```
In [50]: df = pd.DataFrame(passengers, index=pass_id)
```

```
In [51]: print(df.columns)
print(df.info())
print(df.index)

print(df.head(7))
```

```

Index(['age', 'priority'], dtype='object')
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10 entries, 101 to 110
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         10 non-null    int64

```

```

1    priority  10 non-null    object
dtypes: int64(1), object(1)
memory usage: 240.0+ bytes
None
Int64Index([101, 102, 103, 104, 105, 106, 107, 108, 109, 110], dtype='int64')
   age priority
101   23        y
102   25        y
103   78        n
104   12        n
105   56        n
106   33        y
107   67        y

```

3. Select the column of age

```
In [52]: df.age
```

```

Out[52]: 101    23
         102    25
         103    78
         104    12
         105    56
         106    33
         107    67
         108    78
         109    34
         110    64
Name: age, dtype: int64

```

4. Select the rows with id 105, 106 and 109

```
In [53]: df.loc[[105, 106, 109]]
```

```

Out[53]:   age priority
105    56         n
106    33         y
109    34         n

```

5. Get the rows of passengers that are older than 40

```
In [54]: df.loc[df['age'] > 40]
```

```

Out[54]:   age priority
103    78         n
105    56         n
107    67         y
108    78         n
110    64         y

```

6. Assume you get additional passengers from a different source (df2). Concatenate the two DataFrames and name the new DataFrame

df3

```
In [55]: passengers2 = {"age": [54, 65, 12],  
                      "priority": ['y', 'y', 'n']}  
pass_id=[201, 202, 203]
```

```
In [56]: df2 = pd.DataFrame(passengers2, index=pass_id)
```

```
In [57]: df3 = pd.concat([df, df2])  
df3.tail()
```

```
Out[57]:
```

	age	priority
109	34	n
110	64	y
201	54	y
202	65	y
203	12	n

7. Assume that for some passengers we have also some data for the flight they booked. Merge the two DataFrames (df3 and df_flights) including all the rows (outer) (assign the result to a new DataFrame df_merge).

```
In [58]: pass_id = [101, 102, 202, 205, 210, 211]  
flights = ['KL211', 'HV543', 'FR3090', 'KL4345', 'KL4345', 'FR3090']  
df_flights = pd.DataFrame({'flight': flights}, index=pass_id)
```

```
In [59]: print(df3)  
print(df_flights)  
df_merge = pd.merge(df3, df_flights, right_index=True,  
                    left_index=True, how='outer')  
print(df_merge)
```

	age	priority
101	23	y
102	25	y
103	78	n
104	12	n
105	56	n
106	33	y
107	67	y
108	78	n
109	34	n
110	64	y
201	54	y
202	65	y
203	12	n

	flight
101	KL211
102	HV543
202	FR3090
205	KL4345
210	KL4345
211	FR3090

	age	priority	flight
101	23.0	y	KL211
102	25.0	y	HV543
103	78.0	n	NaN
104	12.0	n	NaN
105	56.0	n	NaN
106	33.0	y	NaN
107	67.0	y	NaN
108	78.0	n	NaN
109	34.0	n	NaN
110	64.0	y	NaN
201	54.0	y	NaN
202	65.0	y	FR3090
203	12.0	n	NaN
205	NaN	NaN	KL4345
210	NaN	NaN	KL4345
211	NaN	NaN	FR3090

8. Now you have information on flights delayed. Merge the two DataFrames (delayed with df_merge). Make sure you keep the passengers id

```
In [60]: delayed = pd.DataFrame({'flight': ['KL211','HV543', 'FR3990',
                                             'KL4345','KL4335', 'FR3090'],
                                'delay': ['y', 'n', 'y', 'n', 'y', 'y']})

df_merge_delayed = pd.merge(df_merge.reset_index(), delayed, how = 'left').set_index('index')

df_merge_delayed
```

```
Out[60]:
```

	age	priority	flight	delay
index				
101	23.0	y	KL211	y
102	25.0	y	HV543	n
103	78.0	n	NaN	NaN
104	12.0	n	NaN	NaN
105	56.0	n	NaN	NaN
106	33.0	y	NaN	NaN
107	67.0	y	NaN	NaN
108	78.0	n	NaN	NaN
109	34.0	n	NaN	NaN
110	64.0	y	NaN	NaN
201	54.0	y	NaN	NaN
202	65.0	y	FR3090	y
203	12.0	n	NaN	NaN
205	NaN	NaN	KL4345	n
210	NaN	NaN	KL4345	n
211	NaN	NaN	FR3090	y

9. Set the priority to 'n' for the passenger with id 205 and 210

```
In [61]: df_merge_delayed.loc[[205, 210], ['priority']] = 'n'
```

```
In [62]: df_merge_delayed.loc[[205, 210]]
```

Out[62]:

	age	priority	flight	delay
--	-----	----------	--------	-------

index				
205	NaN	n	KL4345	n
210	NaN	n	KL4345	n

```
In [ ]:
```